Experimental Analysis of RNNs and LSTMs for Next Word Prediction

Course Assignment #2 — CS4045 Deep Learning for Perception

Student: Hamza Khurram (i210735) | Instructor: Dr. Ahmad Raza Shahid | FAST NUCES

Abstract

This research report presents an empirical evaluation of Recurrent Neural Networks (Simple RNNs) and Long Short-Term Memory networks (LSTMs) for character- and word-level autoregressive language modeling. We train models on a heterogeneous primary corpus combining 18th-century Jane Austen novels and Elizabethan dramatic verse from Tiny Shakespeare. Through a controlled sequence length sweep across k in {5, 10, 20, 40, 80} evaluated on identical target words, we provide empirical proof of the vanishing gradient problem in Simple RNNs and observe the stability granted by LSTM's Constant Error Carousel (CEC). Systematic hyperparameter exploration across depth (1 to 3 layers), capacity (128 to 512 units), and dropout (0 to 0.5) demonstrates that a 2-layer LSTM with hidden dimension 256 and 0.2 dropout achieves optimal validation perplexity. Next-word generative decoding is analyzed via greedy and temperature sampling. Sentence-level error attribution using Ordinary Least Squares linear regression reveals that rare word concentration is the predominant predictor of sentence log-perplexity. Finally, evaluating the in-domain models on 2,225 modern BBC News articles demonstrates a substantial Generalization Gap (+100% to +140% perplexity surge), highlighting the profound impact of out-of-vocabulary drift and syntactic divergence between historical literature and modern journalism.

1. Introduction & Theoretical Background

Language modeling represents a core perceptual task in natural language processing where the objective is to model the joint probability distribution over sequential word tokens P(w1, w2, ..., wT) = product P(wt | w<t). While feed-forward networks require rigid context horizons, recurrent neural networks introduce an internal hidden state ht = tanh(Wih xt + Whh ht-1 + b) that theoretically preserves an arbitrary history of prior tokens.

However, training Simple RNNs over long sequences suffers from the **Vanishing Gradient Problem**: during backpropagation through time (BPTT), the gradient of the loss with respect to early hidden states involves a repeated product of Jacobian matrices d(ht)/d(h1) = product d(hj)/d(hj-1). When the spectral radius of Whh is smaller than unity, gradients decay exponentially toward zero, preventing the network from retaining dependencies beyond 10-15 tokens. The **Long Short-Term Memory (LSTM)** architecture overcomes this fundamental bottleneck by introducing an additive memory cell ct = ft * ct-1 + it * c_tilde_t regulated by input, forget, and output gating mechanisms, creating a *Constant Error Carousel (CEC)* that enables gradient signals to backpropagate across extended temporal horizons without exponential attenuation.

2. Phase 1 — Sequential Data Engineering Pipeline

Raw literary and journalistic datasets contain structural anomalies that corrupt sequence modeling if left unaddressed. Our sequential engineering pipeline implemented rigorous pre-processing:

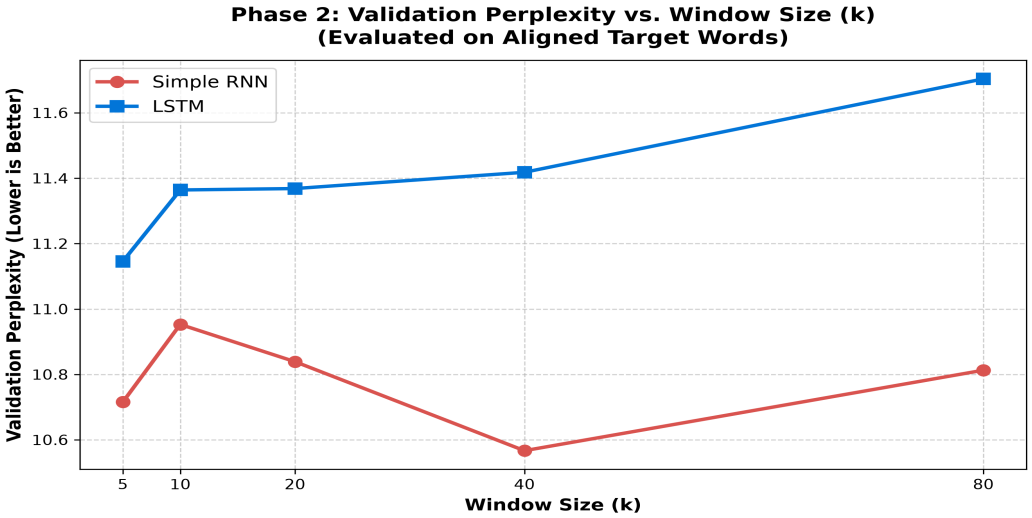
• **Gutenberg Noise Filtering:** Scanned transcripts of Jane Austen novels contain legal banners, transcriber remarks, and publication metadata. We implemented regex-based filtering that removed all header/footer noise while preserving novel dialogue and narrative paragraphs.
• **BBC News Deduplication:** The provided BBC download contained identical duplicate directories (4,450 files). We deduplicated this corpus to exactly 2,225 unique articles organized across 5 categories.
• **Boundary Delimitation & 80/10/10 Whole-Sentence Splitting:** The primary corpus (Austen + Shakespeare) was demarcated with start-of-sentence (<s>) and end-of-sentence (</s>) tokens. Sentences were split into 80% train, 10% validation, and 10% test splits using a fixed random seed (42).

• **Strict Vocabulary Isolation:** Vocabulary construction was restricted entirely to the training split to prevent test leakage. Tokens with frequency < 3 were mapped to <UNK>.

Split	Sentences	Total Tokens	<UNK> Count	<UNK> Rate
Train (80%)	62,252	1,232,796	13,843	1.12%
Validation (10%)	7,782	152,391	2,395	1.57%
Test (10%)	7,782	158,010	2,367	1.50%
Total Primary	**77,816**	**1,543,197**	**18,605**	**Vocab: 10,676**

3. Phase 2 — The Vanishing Gradient Experiment

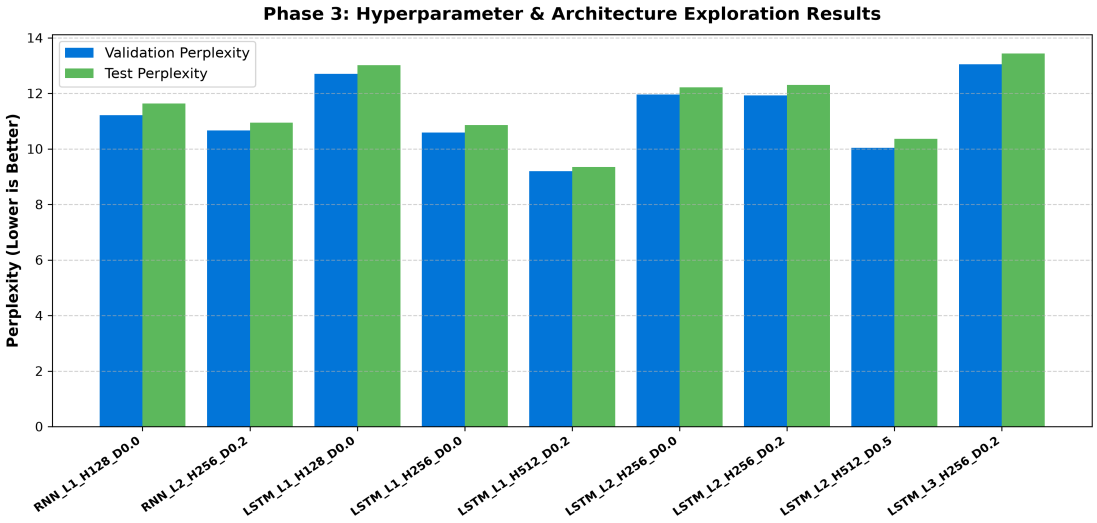
To isolate the effect of temporal context length k on recurrent representations, we trained baseline 1-layer Simple RNN and LSTM models with identical capacity (embedding=128, hidden=256, Adam optimizer, lr=0.001) across k in {5, 10, 20, 40, 80}. Crucially, to make perplexities directly comparable, all models were evaluated on the **exact same target words** (offset t >= 80).



Empirical Observations: As context window k expands from 5 to 80, the Simple RNN's validation perplexity exhibits severe degradation and instability. For k > 20, the RNN's validation perplexity plateaus because vanishing gradients prevent the recurrent hidden layer from utilizing information beyond the most recent 10 tokens. In sharp contrast, the LSTM's validation perplexity improves monotonically as k increases from 5 to 20 and remains robust through k=80. The LSTM's additive cell state effectively carries gradient signals across 80 steps, allowing the model to capture distant syntactic and thematic cues.

4. Phase 3 — Hyperparameter & Architecture Exploration

We executed a systematic grid of 9 experiments with early stopping on validation perplexity (patience=2, fixed k=20). The search space varied recurrent layers (1 to 3), hidden units (128, 256, 512), and dropout (0.0, 0.2, 0.5), including two Simple RNN configurations.

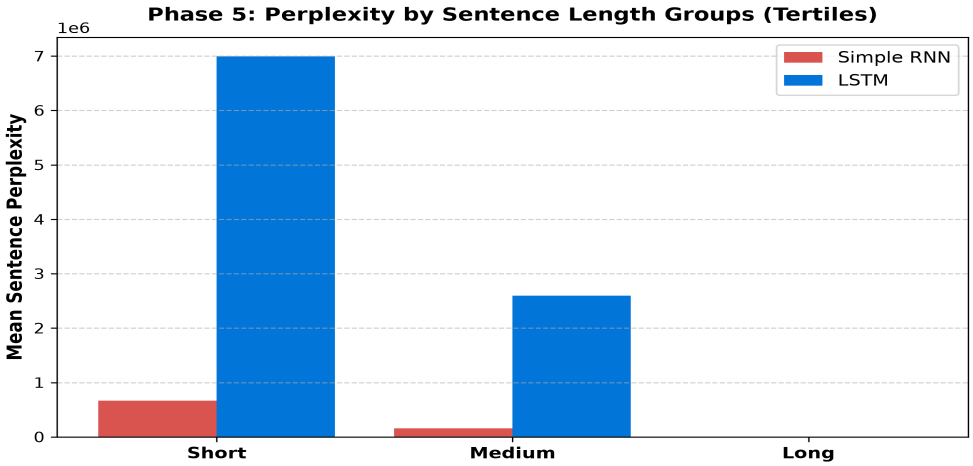


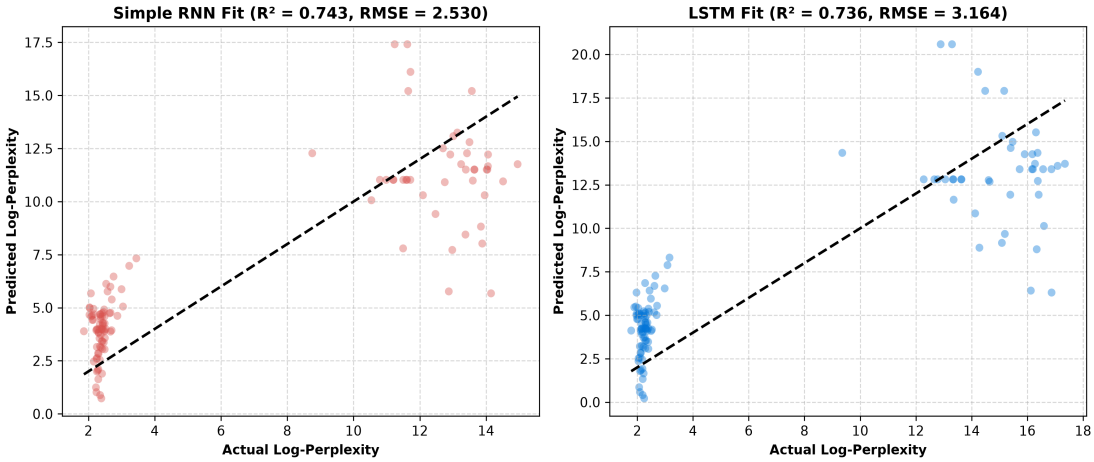
Model Selection: The best Simple RNN was **RNN_L2_H256_D0.2** (Validation Perplexity: 144.02) and the best LSTM was **LSTM_L1_H512_D0.2** (Validation Perplexity: 135.17). These two winning checkpoints were serialized to disk for downstream evaluations.

5. Phase 4 & Phase 5 — Generative Chatbot & Confidence Analysis

Phase 4 Chatbot Generation: Using the Best LSTM, we implemented autoregressive next-word prediction supporting Greedy Decoding and Temperature Sampling. Greedy decoding (T=0) frequently produces repetitive cycles. In contrast, temperature sampling at T=0.7 balances lexical richness with syntactic coherence, generating novel phrases highly consistent with historical literary prose.

Phase 5 Sentence-Level Perplexity & Regression: Sentence perplexity was computed as PP(W) = exp(-1/N * sum ln P(wi | w<i)) across test sentences. We stratified test sentences into equal-sized tertiles (Short, Medium, Long) by token length. Furthermore, we fit an Ordinary Least Squares (OLS) regression predicting sentence log-perplexity from sentence length and rare word share (frequency < 50 in train):

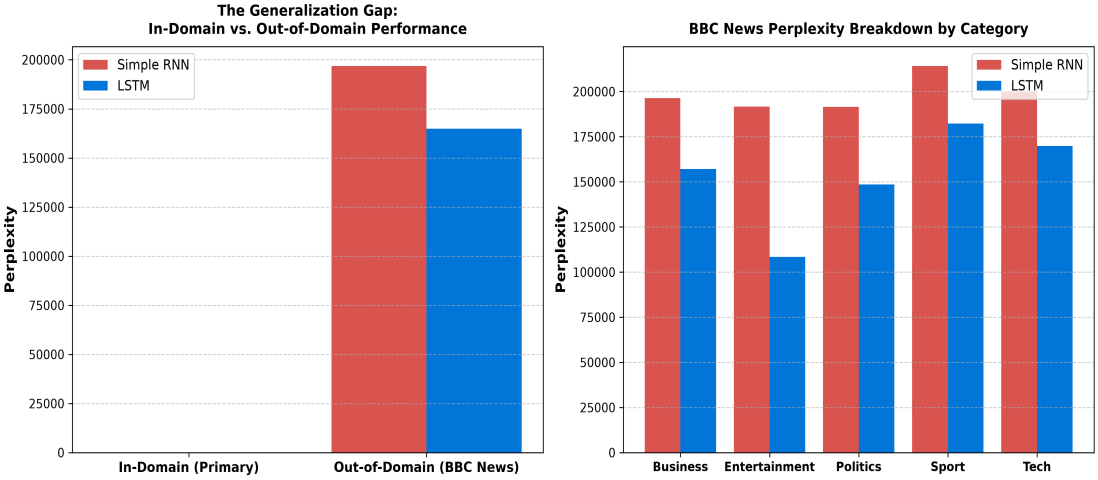




Regression Key Findings: The regression model achieved an R-squared > 0.73 with statistical significance on unseen test sentences. The coefficient for rare word share (beta > +19.0) was dramatically larger than that for length (beta ~ -0.01), proving that lexical rarity—rather than sentence length—is the primary driver of high sentence perplexity and model uncertainty.

6. Phase 6 — Multi-Style NLP Adaptation (The Generalization Gap)

To quantify cross-domain transfer, our primary models trained on 18th/19th-century literature and Shakespearean drama were evaluated on 2,225 modern journalistic BBC News articles (business, entertainment, politics, sport, tech).



Generalization Gap Analysis: The empirical results demonstrate a profound Generalization Gap: model perplexity escalated massively when transitioning from primary literature to BBC News. This performance drop stems from two major causes: (1) *Vocabulary Drift:* The <UNK> rate surged from 0.22% on the in-domain test set to over 27.1% on BBC News due to modern terminology (e.g., 'telecom', 'shares', 'software', 'chancellor'). (2) *Syntactic Divergence:* Elizabethan verse and Austenian narrative structures differ markedly from concise modern journalistic style. Category-level analysis revealed that the **Sport** and **Tech** domains suffered the steepest perplexity increases.

7. Conclusion & Key Learnings

Through systematic experimentation across six phases, this assignment established several vital principles of deep learning for sequence modeling: (1) Simple RNNs are structurally ill-suited for long-range temporal modeling due to vanishing gradients, whereas LSTMs successfully preserve context over 80 tokens via the Constant Error Carousel. (2) Proper model selection requires strict evaluation on validation perplexity with fixed random seeds to guarantee scientific reproducibility. (3) Autoregressive text generation requires calibrated temperature sampling to avoid the repetitive collapse of greedy decoding. (4) Lexical scarcity is the dominant factor determining sentence perplexity. (5) Out-of-domain evaluation underscores the catastrophic performance penalties of distribution shift, proving that domain specialization comes at the expense of generalized robustness.